

1. Google File System (GFS)

Google File System (GFS) is a **scalable distributed file system** developed by Google to handle large data sets using commodity hardware.

Key Characteristics:

- **Fault-tolerant:** Designed to work efficiently even when hardware fails.
- **High throughput:** Optimized for batch processing and large streaming reads.
- **Append-mostly operations:** Supports append-heavy workloads instead of random writes.
- **Scalability:** Easily scales to thousands of nodes and petabytes of data.

Architecture Components:

- **Master Server:**
 - Manages **metadata** (file system tree, chunk locations, access control).
 - Controls **file-to-chunk** mapping and chunk replication.
- **Chunk Servers:**
 - Store actual **data chunks** (default size 64 MB).
 - Handle **read/write** requests from clients.
 - Each chunk is replicated across **multiple chunkservers** (default: 3 replicas).
- **Client:**
 - Requests metadata from **Master** and data directly from **Chunk Servers**.

Working:

- When a file is created, it's split into **chunks** and stored across chunkservers.
- Clients read/write by first contacting the **master** for metadata and then interacting with chunkservers.
- Data consistency is maintained using **version numbers** and **lease mechanisms**.

Fault Tolerance:

- Replication ensures data availability.
- Master periodically **logs and checkpoints metadata**.
- Clients automatically retry on failure.

Limitations:


- Single Master can be a **bottleneck**, though mitigated using **shadow masters** and **checkpointing**.
-

DYNAMO DB

Amazon's relational databases (SQL) struggled with **scale, cost, and complexity** around 2004.

Response:

- In 2007, Amazon released the **Dynamo** paper (core idea of a distributed key-value store).
- In 2012, AWS launched **DynamoDB**, a managed, serverless version of it.



| Feature | SQL (Relational DB) | DynamoDB (NoSQL) | Explanation / Keyword Meaning |
|-------------------|--------------------------|--------------------------------------|--|
| Data Structure | Normalized | Denormalized | <i>Normalized:</i> Data split into multiple tables to avoid redundancy (uses joins). <i>Denormalized:</i> Data kept together for fast access, even with redundancy. |
| Scaling | Vertical Scaling | Horizontal Scaling | <i>Vertical:</i> Add more CPU/RAM to a single server. <i>Horizontal:</i> Add more servers (sharding). |
| Query Flexibility | Ad hoc queries supported | Requires pre-defined access patterns | <i>Ad hoc:</i> Run flexible queries any time (e.g., SQL <code>SELECT * WHERE ...</code>). <i>Pre-defined:</i> Design data upfront for specific queries. |
| Use Case Fit | Good for OLAP | Built for OLTP | <i>OLAP:</i> Analytics-heavy systems (e.g., dashboards). <i>OLTP:</i> High-speed transactional systems (e.g., shopping cart). |

Why not just use SQL everywhere?

→ SQL breaks when traffic, data size, or app complexity explodes.

Scaling Mechanism in DynamoDB

Traditional SQL Scaling (Vertical Scaling):

- In SQL databases, scaling involves upgrading a server's resources (CPU, RAM, disk), which can lead to bottlenecks, downtime, and cost inefficiencies.
- This approach becomes less effective for large-scale, high-traffic applications.

NoSQL Scaling (Horizontal Scaling):

- NoSQL databases like DynamoDB scale horizontally by distributing data across multiple nodes (servers).

- This method is more efficient and handles large, dynamic workloads without bottlenecks.

DynamoDB's Scaling Mechanism:

1. Workload Management:

- Scaling is based on data volume and read/write units (RUs and WUs).
- Each partition can handle up to 1,000 WUs, 3,000 RUs, and 10GB of data.

2. Automatic Incremental Scaling:

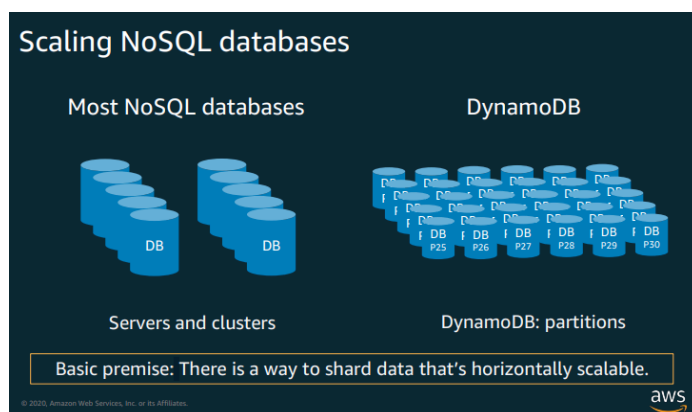
- DynamoDB automatically adjusts the number of partitions based on data access patterns and volume.
- Users interact with tables, while partitioning is handled behind the scenes.

3. Partitioning:

- DynamoDB splits tables into partitions, each assigned to different servers for load distribution.
- This partitioning is transparent to users, who only manage tables.

Key Takeaway:

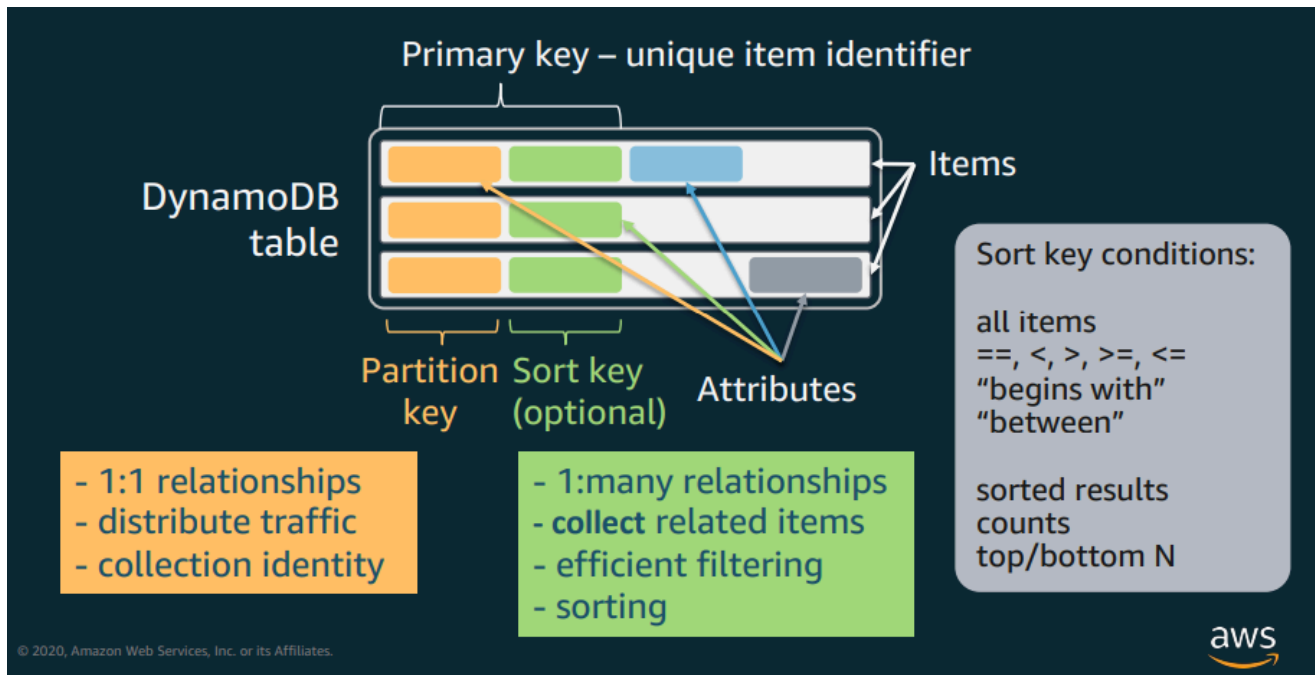
DynamoDB provides automatic horizontal scaling using a partition-based model. This allows it to efficiently handle high traffic and large data volumes without manual intervention.



While most NoSQL systems rely on scaling clusters, **DynamoDB scales horizontally by sharding data into partitions** — making it easier to scale with **high availability, low latency, and automatic load balancing**.

How does it handle spiky loads or high traffic on a single key?

→ Adaptive Capacity & DAX (DynamoDB Accelerator)



DynamoDB Table Structure

- A DynamoDB table stores **items** (similar to rows in a relational database).
- Each item contains **attributes** (similar to columns).

1. Primary Key

- The **primary key** uniquely identifies each item.
- It can be:
 - **Partition Key only** (Simple Primary Key)
 - **Partition Key + Sort Key** (Composite Primary Key)

Purpose

- **Partition Key:**
 - Used to determine the partition where data is stored.
 - Best for one-to-one relationships.
- **Sort Key:**
 - Helps store and retrieve related data.

Partition key

- A good sharding (partitioning) scheme affords even distribution of both data and workload as they grow
- Key concept: **partition key** as the dimension of scalability
 - **Distribute traffic and data** across partitions – horizontal scaling
- Ideal scaling conditions:
 - The partition key is from a high cardinality set (that grows)
 - Requests are evenly spread over the key space
 - Requests are evenly spread over time

- Enables querying within a partition (e.g., time range, top N).

Key Conditions

- Query operations can use conditions like `=`, `<`, `>`, `BETWEEN`, `BEGINS_WITH`, etc., on sort keys.

2. Global Secondary Index (GSI)

- A GSI provides an **alternate way to query data** using different attributes than the primary key.
- It has its own **Partition Key and optional Sort Key**.
- DynamoDB maintains the GSI separately from the main table.

Key Concepts

- **Online Indexing:** The index is updated as new data is written.
- **Projected Attributes:**
 - You can choose to include specific attributes from the main table.
 - Three types:
 - `KEYS_ONLY`: only partition and sort keys
 - `INCLUDE`: specific non-key attributes
 - `ALL`: all attributes
- You can define **up to 20 GSIs per table**.
- **Capacity is provisioned separately** from the main table. [Chatgpt Example](#)

| Feature | Primary Index | Global Secondary Index (GSI) |
|-----------------|------------------------------|---------------------------------|
| Defined when | Table is created | Created optionally anytime |
| Based on | Table's Partition & Sort Key | Custom attributes you choose |
| Can be changed? | No | Yes, can be added later |
| Capacity | Uses table's capacity | Has its own capacity settings |
| Use case | Default, most frequent query | Alternate access/query patterns |

Sharding / Partitioning

- DynamoDB automatically shards data across multiple partitions for scalability and performance.

How It Works

- Each item's **Partition Key** is hashed.
- The hash output determines which **partition** the item will go to.
- Related items (with the same partition key) are stored together.

Benefits

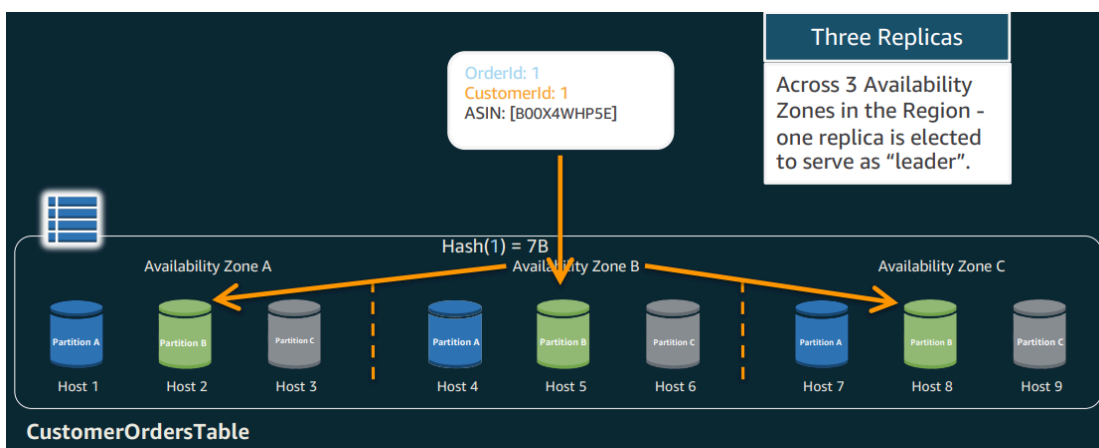
- Prevents uneven load (avoids “hot partitions”).
- Ensures data is spread across nodes for parallelism.
- Supports millions of reads/writes per second.

A View from a Different Angle – Replication and Availability

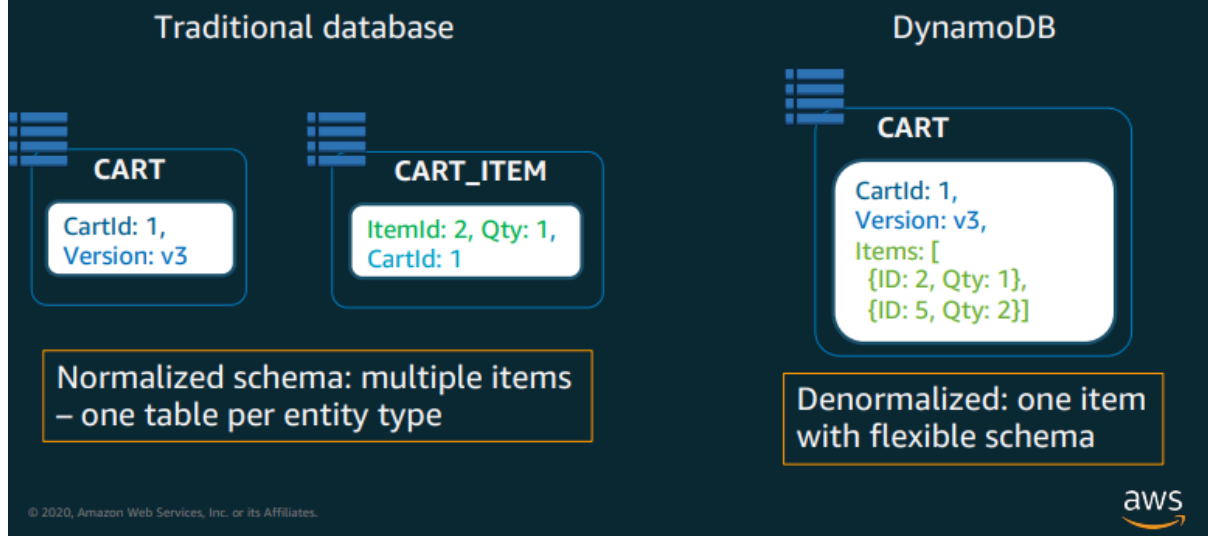
- Each DynamoDB partition is **replicated across 3 Availability Zones (AZs)** within an AWS Region.

Key Points

- **Three replicas** per partition ensure fault tolerance.
- One replica acts as the **leader**, handling all writes.
- The other two are **followers**, available for reads or takeover if needed.
- This architecture ensures:
 - **High Availability**
 - **Durability**
 - **Automatic failover** in case of an AZ outage



Denormalization



Denormalization is the process of combining related data into a single structure instead of splitting it across multiple related tables.

Comparison with Traditional Databases

- Traditional databases use normalization with multiple related tables (e.g., **Cart** and **Cart_Item**)
- Enforce strict schemas and use joins

DynamoDB Approach

- Uses **denormalization**: stores all related data in a single record
- Schema-less structure allows faster lookups

Traditional SQL Queries:

```
sql
SELECT * FROM cart WHERE id=1;
SELECT * FROM cart_items WHERE cart_id=1;
```

Equivalent DynamoDB Record:

```
json
{
  "CartID": 1,
  "Version": 1,
  "Items": [
    {"ID": "X", "Qty": 1},
    {"ID": "Y", "Qty": 2}
  ]
}
```

- All data is fetched and updated together

Ensuring data consistency on updates

Example: add/remove shopping cart items

```
1. get cart => vread = Version
2. update cart:
  IF Version = vread
    add/remove cart items
    ++Version
  ELSE go back to Step 1.
```

Use **ConditionExpression** in DynamoDB

Optimistic concurrency control

© 2020, Amazon Web Services, Inc. or its Affiliates.



Problem

Concurrent updates (e.g., multiple users updating a cart) can lead to **race conditions**.

Solution: Optimistic Concurrency Control (OCC)

- Read the current version (e.g., version = 3)
- Update only if the version hasn't changed
- If the version has changed, retry

DynamoDB Implementation

json

Copy

Edit

```
"ConditionExpression": "version = :v",
"ExpressionAttributeValues": {
  ":v": 3
}
```

- Update will fail if the version doesn't match

Updating cart: data consistency using OCC

1. Get the cart: **GetItem**

```
{ "TableName": "Cart",
  "Key": {"CartId": {"N": "2"}}
}
```



2. Update the cart: conditional **PutItem** (or **UpdateItem**)

```
{ "TableName": "Cart",
  "Item": {
    "CartID": {"N": "2"},
    "Version": {"N": "4"},
    "CartItems": {...}
  },
  "ConditionExpression": "Version = :ver",
  "ExpressionAttributeValues": {":ver": {"N": "3"}}
}
```

- ✓ Use conditions to implement optimistic concurrency control ensuring data consistency
- ✓ Single-item operations are ACID
- ✓ GetItem call can be eventually consistent

© 2020, Amazon Web Services, Inc. or its Affiliates.



Yes, you **can have strongly consistent read-after-write and concurrency control** with DynamoDB — here's the justification:

1. Strongly Consistent Read-after-Write

By default, DynamoDB provides **eventual consistency**, meaning you might get stale data right after a write.

But if you need **immediate accuracy** (like verifying a recent update), **DynamoDB supports *strongly consistent reads*** as an option.

How?

- When you make a read request, you can set:

json

Copy Edit

```
"ConsistentRead": true
```

- This ensures the read always reflects **the latest successful write** (from the same region).

Example Use Case:

- After updating a cart, you immediately fetch the cart again and want to ensure you get the updated version — strongly consistent read guarantees that.

2. Concurrency Control via Optimistic Locking (OCC)

DynamoDB allows **safe concurrent writes** using **Optimistic Concurrency Control**, which avoids race conditions.

How it works:

- You store a **version** field in your item.
- When updating, you include a **ConditionExpression** that checks if the version hasn't changed.

Example:

json

Copy Edit

```
"ConditionExpression": "version = :oldv",  
"ExpressionAttributeValues": {  
  ":oldv": 3  
}
```

- If someone else updated the item and the version is now 4, your update **fails gracefully**, preventing conflicts.
- You can then **retry** by re-fetching the latest item and attempting the update again.

| Feature | Support in DynamoDB | Purpose |
|---------------------------|---|-----------------------------------|
| Strongly Consistent Read | Yes (<code>ConsistentRead: true</code>) | Ensures accurate read after write |
| Concurrency Control (OCC) | Yes (using <code>ConditionExpression</code>) | Prevents conflicting writes |

✓ By Default: Eventually Consistent

- When you **read data** after a write, you **might not see the latest value immediately**.
- This is called **eventual consistency** — the system guarantees that **all replicas will be updated eventually**, but **not instantly**.

🧠 Why Use Eventual Consistency?

- It's **faster** and more **scalable**, especially for global-scale applications.
- It reduces the load on the system and helps maintain high availability.

🔒 But... Strong Consistency is Also Available!

If your use case needs **immediate consistency**, DynamoDB lets you **explicitly request it**:

json

📄 Copy

✎ Edit

```
"ConsistentRead": true
```

This makes the read **strongly consistent**, meaning:

- You'll **always get the latest committed data** from a successful write.
- But this works **only within the same AWS Region** (not across regions).

Yes, you can model complex data relationships with DynamoDB

DynamoDB Data Modeling Patterns

1. One-to-One (Key-Value Lookup):

- Use **Partition Key** (e.g., `UserId` or `Email`)
- Access via `GetItem` / `BatchGetItem`
- Example: Get user details by email.

2. One-to-Many (Parent-Child):

- Use **Partition Key + Sort Key**
- Access via `Query`
- Example: Get all device readings between two timestamps.

3. Many-to-Many:

- Use a table and a **GSI with PK and SK swapped**
- Access via `Query`
- Example: Find all games for a user or all users for a game.

Challenges

- Growing datasets
- Variable or spiky throughput
- Imbalanced workloads (hot partitions)
- Poor traffic distribution across shards

♦ Solutions

1. Time-To-Live (TTL):

- Automatically delete expired data to reduce storage.

2. Provisioned Mode + Auto Scaling:

- Automatically adjusts capacity based on usage patterns.

3. On-Demand Mode:


- Ideal for unpredictable or spiky workloads. No capacity planning needed.

4. Adaptive Capacity:

- DynamoDB shifts capacity to "hot" partitions to handle uneven workloads.

5. DynamoDB Accelerator (DAX):

- In-memory caching layer for **read-heavy** applications
- Fully managed, fault-tolerant, multi-AZ
- Compatible with DynamoDB API (no code changes needed)
- Supports **write-through caching**



Integrating DynamoDB into your data flow

1. Complex Queries and Analytics

Problem

Amazon DynamoDB is a high-performance NoSQL database. It is optimized for key-value access and simple queries, but it is **not designed for complex analytics or full-text search**. For these use cases, you may need other services.

Solution

To handle complex analytics and queries:

- **Use services best suited for those operations**, such as:
 - **Amazon Athena** – for interactive SQL queries on data stored in S3.

- **Amazon Redshift** – a data warehouse optimized for analytical queries.
- **Amazon Elasticsearch Service (now OpenSearch Service)** – for full-text search and log analytics.

Data Sync Using DynamoDB Streams

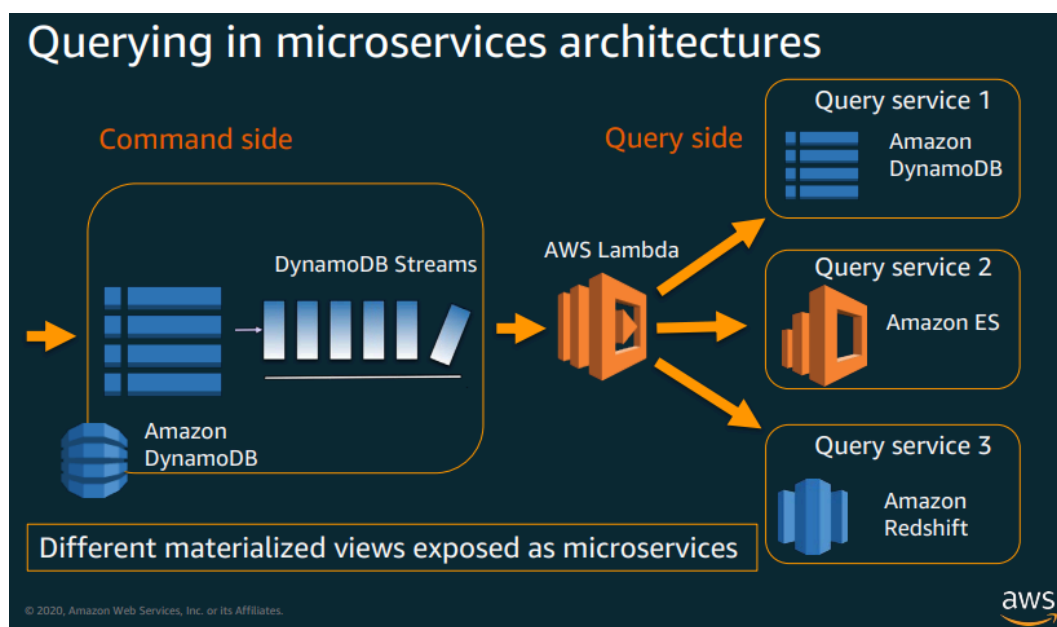
To **ensure data consistency** across these services and maintain up-to-date information:

- **DynamoDB Streams** captures real-time changes (insert/update/delete) in a DynamoDB table.
- These changes can be **consumed by an AWS Lambda function**.
 - AWS Lambda is a serverless compute service that can execute code in response to DynamoDB stream events.
 - It allows for **end-to-end serverless architecture** with no manual intervention or servers to manage.

Workflow

- Data is written to the **DynamoDB table**.
- Changes are pushed to **DynamoDB Streams**.
- **AWS Lambda** consumes these changes and processes them (e.g., updates an analytics system or another database).
- Final analytics results can be stored or queried from a more appropriate service.

2. Querying in Microservices Architectures



Architecture Breakdown

This section shows how **DynamoDB Streams + AWS Lambda** can be used to enable **multiple microservices to work with different views of the same data**.

Separation of Concerns

- **Command Side** (Write Operations):
 - All writes (create/update/delete) are performed on the **Amazon DynamoDB** table.
 - These changes are captured by **DynamoDB Streams** in real time.
- **Query Side** (Read Operations):
 - Using **AWS Lambda**, changes from the stream are processed and sent to different **query services**.
 - Each query service has its own database or system optimized for its specific query requirements.

Examples of Query Services

- **Query Service 1** – continues to use **Amazon DynamoDB** if fast key-based lookups are sufficient.
- **Query Service 2** – sends data to **Amazon Elasticsearch Service (ES)** for full-text search capabilities.
- **Query Service 3** – sends data to **Amazon Redshift** for complex analytical queries.

Materialized Views as Microservices

Each query service can be thought of as a **materialized view** – a precomputed dataset tailored for a specific use case or query pattern.

- These views are maintained **automatically** through the stream + lambda mechanism.
- This pattern is commonly referred to as **CQRS (Command Query Responsibility Segregation)**.

Summary: Key Takeaways

1. **DynamoDB Streams + AWS Lambda** enables **real-time data propagation** from your core DynamoDB database to other specialized services.
2. **Decouple writes and reads** – you write once to DynamoDB and create many **query-optimized views** in different services.
3. Supports **serverless architecture** – no infrastructure management needed.
4. Enables **scalable microservices design** where different parts of your application (or different teams) can independently scale and optimize their data usage.

Polyglot Persistence – in simple terms:

Definition:

It means **using different types of databases** (SQL, NoSQL, graph, etc.) for different parts of an application — **"the right tool for the right job."**

♦ Why it's used?

Because **no single database** is perfect for every use case. For example:

- **DynamoDB** for high-scale key-value or document data
 - **RDS/MySQL** for relational, transactional data
 - **Elasticsearch** for full-text search
 - **Redis** for caching
 - **Neo4j** (graph DB) for relationship-heavy data
-

♦ Example:

An **e-commerce site** might use:

- **DynamoDB** for user sessions (fast, scalable)
- **MySQL** for orders and payments (ACID transactions)
- **Redis** for quick product lookups
- **Elasticsearch** for search functionality